

REPORT

Building EDF Scheduler based on FreeRTOS Kernel.

BY: MOHAMMED KHALID

Entrance

FreeRTOS is one of the real time operating systems that uses the fixed priority schedulers to handle tasks executions within a computer machine, that is, every task in the system is pre-assigned a priority value while the system is being designed to represent its weightiness among others tasks.

In the fixed priority topology, it's up to designer to consider every task's deadline, and to choose carefully the priorities in order for the system to be schedulable and all tasks be able to catch their deadlines.

Although this might appear to put extra load on a designer's shoulder, This topology is doing well in many cases where task's deadline is not the most critical part of the system. In others systems, like safety-critical ones, tasks can be event triggered, and catching their deadline is the most important criteria to achieve. That's why the fixed priority scheduler has some limitations at those kind of systems.

Although, FreeRTOS is offering nice APIs allows a programmer from managing the priorities at runtime and enabling a system to act dynamically, still there is too much headache to relate a task deadline with a dynamic priority.

That's why other typologies in scheduling like EDF, or Earliest Deadline First, is exist. EDF scheduler automatically relates deadlines and prioritization, where at any point of time the task being executed has nearest deadline between all other tasks, and no other task can preempt it unless it has has earliest deadline than it.

Here we present an algorithm to built such scheduler based on the FreeRTOS Kernel implementation and make it compatible with the original fixed priority scheduler to make it easy for switching between them.

The algorithm assigns priorities to tasks in a simple way: the priority of a task is inversely proportional to its absolute deadline; In other words, the highest priority is the one with the earliest deadline. In case of two or more tasks with the same absolute deadline, the highest priority task among them is chosen random.

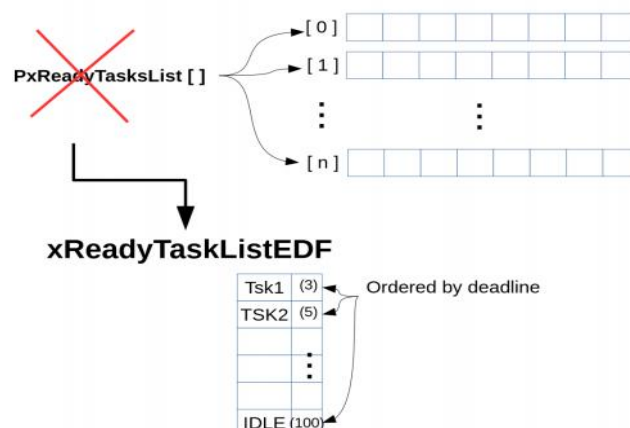
■ **The algorithm is suited to work in an environment where these assumptions applies :-**

1. The requests for all tasks for which hard deadlines exist are periodic, with constant interval between requests.
2. Deadlines consist of run-ability constraints only, i.e. each task must be completed before the next requests for it occurs.
3. The tasks are independent in that requests for a certain task do not depend on the initialization or the completion of requests for other tasks.
4. Run-time for each task is constant for that task and does not vary with time. Run-time refers to the time which is taken by a processor to execute the task without interruption.

Implementation in FreeRTOS Kernel

As said previously, FreeRTOS uses a scheduler based on static priority policy. The aim of this chapter is to describe how to implement an EDF scheduler, using the existing structures that FreeRTOS offers and creating new ones. The general idea is to create a new Ready List able to manage a dynamic task priority behaviour: it will contain tasks ordered

by increasing deadline time, where positions in the list represent the tasks priorities, with the head of the list containing the running task. The rest of FreeRTOS architecture and structures, as the Waiting List and the clock mechanism are maintained with marginal changes.



All changes that will be illustrated refer to tasks.c file, since scheduler structures and methods are contained there. According with the FreeRTOS style guideline, a configuration variable ***configUSE_EDF_SCHEDULER***, is added to the ***FreeRTOSConfig.h*** file. When ***configUSE_EDF_SCHEDULER***, is set to 1, EDF scheduler is used, elsewhere the OS uses the original scheduler.

```

42  * See http://www.freertos.org/a00110.html
43  *-----*/
44
45  #define configUSE_EDF_SCHEDULER 1
46
47  #define configUSE_PREEMPTION 1

```

Then, the new Ready List is declared: ***xReadyTasksListEDF*** is a simple list structure.

```

197  /*-----*/
198  /**
199  * When configUSE_EDF_SCHEDULER is set to 1, EDF scheduler is used, declare the EDF scheduler Ready List.
200  */
201  #if ( configUSE_EDF_SCHEDULER == 1 )
202
203  PRIVILEGED_DATA static List_t xReadyTasksListEDF; /*< Ready tasks ordered by their deadline. */
204
205  #endif
206
207  /*-----*/

```

Then, the ***prvInitialiseTaskLists()*** method, that initialize all the task lists at the creation of the first task, is modified adding the initialization of ***xReadyTasksListEDF***:

```

3678 static void prvInitialiseTaskLists( void )
3679 {
3680
3681     #if ( configUSE_EDF_SCHEDULER == 1 )
3682
3683         vListInitialise( &xReadyTasksListEDF );
3684
3685     #else
3686         UBaseType_t uxPriority;
3687         for( uxPriority = ( UBaseType_t ) 0U; uxPriority < ( UBaseType_t ) configMAX_PRIORITIES; uxPriority++ )
3688         {
3689             vListInitialise( &(amp; pxReadyTasksLists[ uxPriority ] ) );
3690         }
3691     #endif

```

prvAddTaskToReadyList() method that adds a task to the Ready List is then modified as follows:

```

229 /*
230  * Place the task represented by pxTCB into the appropriate ready list for
231  * the task. It is inserted at the end of the list.
232  */
233 #if ( configUSE_EDF_SCHEDULER == 0 )
234     #define prvAddTaskToReadyList( pxTCB ) \
235         traceMOVED_TASK_TO_READY_STATE( pxTCB ); \
236         taskRECORD_READY_PRIORITY( ( pxTCB )->uxPriority ); \
237         listINSERT_END( &( pxReadyTasksLists[ ( pxTCB )->uxPriority ] ), &( ( pxTCB )->xStateListItem ) ); \
238         tracePOST_MOVED_TASK_TO_READY_STATE( pxTCB )
239 #else
240     #define prvAddTaskToReadyList( pxTCB ) \
241         traceMOVED_TASK_TO_READY_STATE( pxTCB ); \
242         vListInsert( &(xReadyTasksListEDF), &( (pxTCB)->xStateListItem ) ) | \
243         tracePOST_MOVED_TASK_TO_READY_STATE( pxTCB )
244 #endif
245
246 /*-----*/
247

```

vListInsert() method is called to insert in ***xReadyTasksListEDF*** the task **TCB** pointer. The item will be inserted into the list in a position determined by its item value ***xGenericListItem*** (descending item value order). So it is assumed that ***xStateListItem*** contains the next task deadline.

The second change introduced refers to the task structure. When a task moves to the Ready List, the knowledge of its next deadline is needed in order to insert it in the correct position. The deadline is calculated as:

$TASK_{deadline} = Tick_{cur} + TASK_{period}$, so every task needs to store its period value. A new variable is added in the ***tskTaskControlBlock*** structure (**TCB**):

```

284 ListItem_t xStateListItem; /*< The list that the state list item of a task is reference from denotes t
285 ListItem_t xEventListItem; /*< Used to reference a task from an event list. */
286 UBaseType_t uxPriority; /*< The priority of the task. 0 is the lowest priority. */
287 StackType_t * pxStack; /*< Points to the start of the stack. */
288 char pcTaskName[ configMAX_TASK_NAME_LEN ]; /*< Descriptive name given to the task when created. Facilitates debugging
289
290 #if ( configUSE_EDF_SCHEDULER == 1 )
291     TickType_t xTaskPeriod; /*< Stores the period in tick of the task. > */
292 #endif
293
294

```

Accordingly, a new initialization task method is created. ***xTaskPeriodicCreate()*** is a modified version of the standard method ***xTaskCreate()***, that receives the task period as additional input parameter and set the ***xTaskPeriod*** variable in the task **TCB** structure.

```

847
848 #if ( configUSE_EDF_SCHEDULER == 1 )
849     BaseType_t xTaskPeriodicCreate( TaskFunction_t pxTaskCode,
850                                     const char * const pcName, /*lint !e971 Unqualified char types are allowed for strin
851                                     const configSTACK_DEPTH_TYPE usStackDepth,
852                                     void * const pvParameters,
853                                     UBaseType_t uxPriority,
854                                     TaskHandle_t * const pxCreatedTask, TickType_t period )
855     {
856         TCB_t * pxNewTCB;
857

```

Before adding the new task to the Ready List by calling ***prvAddTaskToReadyList()***, the task's ***xStateListItem*** is initialized to the value of the next task deadline.

```

926      /*E.C. : initialize the period */
927      pxNewTCB->xTaskPeriod = period;
928
929      prvInitialiseNewTask( pxTaskCode, pcName, ( uint32_t ) usStackDepth, pvParameters, uxPriority, pxCreatedTask, pxNewTCB, NULL );
930      /*E.C. : insert the period value in the generic list item before to add the task in RL: */
931      listSET_LIST_ITEM_VALUE( &( ( pxNewTCB )->xStateListItem ), (pxNewTCB)->xTaskPeriod + xTickCount);
932
933      prvAddTaskToReadyList( pxNewTCB );
934
935      xReturn = pdPASS;

```

The IDLE task management is modified as well. The initialization of the IDLE task happens in the ***vTaskStartScheduler()*** method, that starts the real time kernel tick processing and initialize all the scheduler structures. Since FreeRTOS specifications want a task in execution at every instant, a correct management of the IDLE task is fundamental. With the standard FreeRTOS scheduler, the IDLE task is a simple task initialized at the lowest priority. In this way it would be scheduled only when no other tasks are in the ready state. With the EDF scheduler, the lowest priority behaviour can be simulated by a task having the farrest deadline.

vTaskStartScheduler() method initializes the IDLE task and inserts it into the Ready List. The method is modified as follow:

```

2165  #elseif (configUSE_EDF_SCHEDULER == 1)
2166  {
2167      /* The Idle task is being created with the farrest deadline. */
2168      TickType_t initIDLEPeriod = Init_Idle_Period;
2169
2170      xReturn = xTaskPeriodicCreate( prvIdleTask,
2171                                   configIDLE_TASK_NAME,
2172                                   configMINIMAL_STACK_SIZE,
2173                                   ( void * ) NULL,
2174                                   portPRIVILEGE_BIT, /* In effect ( tskIDLE_PRIORITY | portPRIVILEGE_BIT ), but tskIDLE_PRIORITY is zero. */
2175                                   &xIdleTaskHandle, initIDLEPeriod ); /*lint !e961 MISRA exception, justified as it is not a redundant explicit
2176  }

```

Where ***Init_Idle_Period*** is a user defined macro ..

```

67
68
69  #define Init_Idle_Period      ((unsigned long) 10000)
70
71

```

The IDLE task is initialized with a period of ***initIDLEPeriod*** = 10000. We assume that no task can have a period greater than ***initIDLEPeriod***: in this way, when the IDLE task is

added to the Ready List, it will be at the last position of the list, since its deadline will be greater than any other task ($TASK_{deadline} = Tick_{cur} + TASK_{period}$) with $Tick_{cur} = 0$ and ***IDLE_{period}*** = ***initIDLEPeriod*** greater than any other task period.

Every time IDLE task executes (i.e. no other tasks are in the Ready List), it calls a method that increments its deadline in order to guarantee that IDLE task will remain in the last position of the Ready List.

```

2972
2973     #if (configUSE_EDF_SCHEDULER == 1)
2974         /*< calculate the new task deadline, if EDF scheduler is used>*/
2975         listSET_LIST_ITEM_VALUE( &( ( pxTCB )->xStateListItem ), ( pxTCB->xTaskPeriod + xTickCount);
2976
2977     #endif
2978
2979     /* Place the unblocked task into the appropriate ready
2980      * list. */
2981     prvAddTaskToReadyList( pxTCB );
2982

```

Another change needed involves the switch context mechanism. Every time the running task is suspended, or a suspended task with an higher priority than the running task awakes, a switch context occurs. ***vTaskSwitchContext()*** method is in charge to update the ****pxCurrentTCB*** pointer to the new running task:

```

3225
3226     #if (configUSE_EDF_SCHEDULER == 0)
3227         taskSELECT_HIGHEST_PRIORITY_TASK(); /*lint !e9079 void * is used as this macro is used with timers and co-r
3228     #else
3229         pxCurrentTCB = (TCB_t *) listGET_OWNER_OF_HEAD_ENTRY( &(xReadyTasksListEDF) );
3230     #endif
3231

```

taskSELECT_HIGHEST_PRIORITY_TASK() method is replaced in order to assign to ***pxCurrentTCB*** the task at the first place of the new Ready List.

Last change required to get all the pieces in the new EDF scheduler to work is to update tasks deadlines stored in the ***xStateListItem*** every tick, but actually we just need to update deadline for those tasks whose deadlines outdated, that is the tasks has recently unblocked from the waiting list, again the new deadline will be $TASK_{deadline} = Tick_{cur} + TASK_{period}$.

```

2972
2973     #if (configUSE_EDF_SCHEDULER == 1)
2974         /*< calculate the new task deadline, if EDF scheduler is used>*/
2975         listSET_LIST_ITEM_VALUE( &( ( pxTCB )->xStateListItem ), ( pxTCB->xTaskPeriod + xTickCount);
2976
2977     #endif
2978
2979     /* Place the unblocked task into the appropriate ready
2980      * list. */
2981     prvAddTaskToReadyList( pxTCB );
2982

```

Also we have determine if a context switching is required after this tick or not. If a released task has earliest deadline than the one that being executed then we need a ***Context Switching***, then we compare both of them in order make the decision.


```

2979          /* Place the unblocked task into the appropriate ready
2980             * list. */
2981          prvAddTaskToReadyList( pxTCB );
2982
2983          /* If the EDF scheduler is used and the task being unblocked has a nearest deadline
2984             than the current task deadline, a Context Switch is required*/
2985          #if (configUSE_EDF_SCHEDULER == 1)
2986
2987              if ( listGET_LIST_ITEM_VALUE(&(( pxTCB )->xStateListItem )) <=
2988                  listGET_LIST_ITEM_VALUE(&(( pxCurrentTCB )->xStateListItem )) )
2989              {
2990
2991                  xSwitchRequired = pdTRUE;
2992              }
2993              else
2994              {
2995                  mtCOVERAGE_TEST_MARKER();
2996              }
2997
2998          #endif
2999

```

Testing and Analysis

To test and verify our implementation of the EDF scheduler we aim to create 6 tasks with different deadlines and different execution time.

The six tasks are as follows :-

Task 1: "Button_1_Monitor" , {Periodicity: 50, Deadline: 50}

This task will monitor rising and falling edge on button 1 and send this event to the consumer task.

Task 2: ""Button_2_Monitor"", {Periodicity: 50, Deadline: 50}

This task will monitor rising and falling edge on button 2 and send this event to the consumer task.

Task 3: ""Periodic_Transmitter"", {Periodicity: 100, Deadline: 100}

This task will send periodic string every 100ms to the consumer task.

Task 4: ""Uart_Receiver"", {Periodicity: 20, Deadline: 20}

This is the consumer task which will write on UART any received string from other tasks.

Task 5: "Load_1_Simulation", {Periodicity: 10, Deadline: 10}, with execution time = 5ms.

Task 6: "Load_2_Simulation", {Periodicity: 100, Deadline: 100}, with execution time = 12ms.

And to trace the tasks execution and scheduling at runtime we use the Trace Hooks macros offered by FreeRTOS.

| Macro definition | Description |
|---|---|
| <code>traceTASK_INCREMENT_TICK(xTickCount)</code> | Called during the tick interrupt. |
| <code>traceTASK_SWITCHED_OUT()</code> | Called before a new task is selected to run. At this point <code>pxCurrentTCB</code> contains the handle of the task about to leave the Running state. |
| <code>traceTASK_SWITCHED_IN()</code> | Called after a task has been selected to run. At this point <code>pxCurrentTCB</code> contains the handle of the task about to enter the Running state. |


Using the GPIOs pin to generate signals in tasks going it and out from execution we can analyse our system at runtime..

```

90
91  /*****
92      Trace Hook Macros
93  *****/
94
95  #define traceTASK_INCREMENT_TICK(xTickCount) \
96      GPIO_write(PORT_0, PIN0, PIN_IS_HIGH); \
97      GPIO_write(PORT_0, PIN0, PIN_IS_LOW);
98
99
100 #define traceTASK_SWITCHED_OUT() \
101     if ((int)pxCurrentTCB->pxTaskTag == 1) \
102         GPIO_write(PORT_0, PIN1, PIN_IS_LOW); \
103     if ((int)pxCurrentTCB->pxTaskTag == 2) \
104         GPIO_write(PORT_0, PIN2, PIN_IS_LOW); \
105     if ((int)pxCurrentTCB->pxTaskTag == 3) \
106         GPIO_write(PORT_0, PIN3, PIN_IS_LOW); \
107     if ((int)pxCurrentTCB->pxTaskTag == 4) \
108         GPIO_write(PORT_0, PIN4, PIN_IS_LOW); \
109     if ((int)pxCurrentTCB->pxTaskTag == 5) \
110         GPIO_write(PORT_0, PIN5, PIN_IS_LOW); \
111     if ((int)pxCurrentTCB->pxTaskTag == 6) \
112         GPIO_write(PORT_0, PIN6, PIN_IS_LOW);
113
114
115 #define traceTASK_SWITCHED_IN() \
116     if ((int)pxCurrentTCB->pxTaskTag == 1) \
117         GPIO_write(PORT_0, PIN1, PIN_IS_HIGH); \
118     if ((int)pxCurrentTCB->pxTaskTag == 2) \
119         GPIO_write(PORT_0, PIN2, PIN_IS_HIGH); \
120     if ((int)pxCurrentTCB->pxTaskTag == 3) \
121         GPIO_write(PORT_0, PIN3, PIN_IS_HIGH); \
122     if ((int)pxCurrentTCB->pxTaskTag == 4) \
123         GPIO_write(PORT_0, PIN4, PIN_IS_HIGH); \
124     if ((int)pxCurrentTCB->pxTaskTag == 5) \
125         GPIO_write(PORT_0, PIN5, PIN_IS_HIGH); \
126     if ((int)pxCurrentTCB->pxTaskTag == 6) \
127         GPIO_write(PORT_0, PIN6, PIN_IS_HIGH);

```

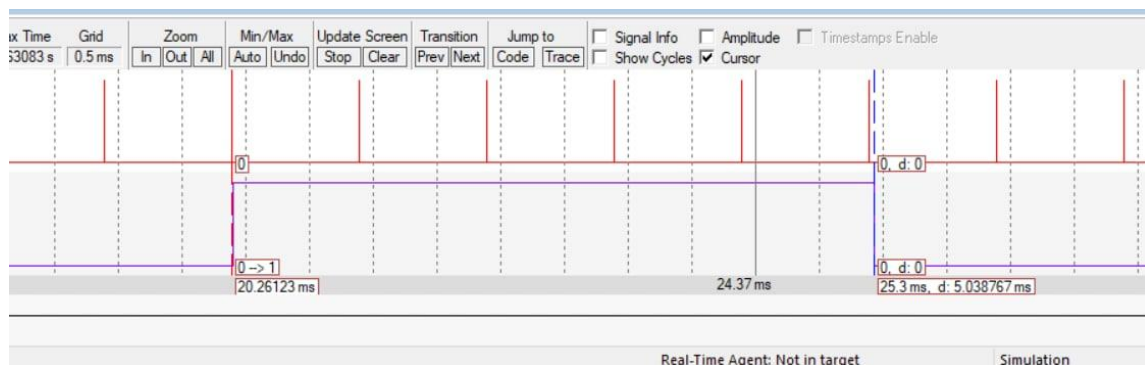
We devote six pins for tracing six task (from PIN1 to PIN6) , and PIN0 for the tick routine.



With using of logic analyzer we can see a real time line contain the six tasks and their switching in execution.

Hence, we can also tune our load simulation tasks to achieve the required execution time.


One has 5ms execution time and the other is 12 ms..



Also to verify our system correctness of scheduling we use the Simso tool to get another accurate Time-line which can be compared with ours.

| Model data | | | | | | | | | |
|------------|---------|-----------|---|----------------|-------------|-------------------------|---------------|-----------|-------------|
| General | | Scheduler | Processors | Tasks | | | | | |
| id | Name | Task type | Abort on miss | Act. Date (ms) | Period (ms) | List of Act. dates (ms) | Deadline (ms) | WCET (ms) | Followed by |
| 1 | TASK T1 | Periodic | <input checked="" type="checkbox"/> Yes | 0 | 50 | - | 50 | 0.012 | |
| 2 | TASK T2 | Periodic | <input checked="" type="checkbox"/> Yes | 0 | 50 | - | 50 | 0.012 | |
| 3 | TASK T3 | Periodic | <input checked="" type="checkbox"/> Yes | 0 | 100 | - | 100 | 0.023 | |
| 4 | TASK T4 | Periodic | <input checked="" type="checkbox"/> Yes | 0 | 20 | - | 20 | 0.016 | |
| 5 | TASK T5 | Periodic | <input checked="" type="checkbox"/> Yes | 0 | 10 | - | 10 | 5 | |
| 6 | TASK T6 | Periodic | <input checked="" type="checkbox"/> Yes | 0 | 100 | - | 100 | 12 | |

we calculate every task execution time in our system and configure Simso with these information to get us this illustration graph.



It's obvious that it seems identical to our Time-Line graph, hence our scheduler behaves correctly!

Another important calculation we have to consider is the system **hyperperiod** and **The CPU load**.

Hyperperiod is the point in where in each the system repeating it self, in other words, all tasks are to be scheduled in the same time. Hyperperiod period is usually The greatest common factor among all the tasks periods.

Here, our GCF is 100ms. And it's clear on the time line that after 100ms and 200ms and off course every 100ms all tasks are came to be scheduled in these moments.

And for the Task.5 to have the least deadline among other 5 tasks it's scheduled to execute first.



Then to Calculate the CPU usage time we have to utilize a general purpose timer with frequency greater than the Tick Timer, in this purpose we use the internal Timer1 in our MCU.

And to facilitate the gathering of information we use the ready FreeRTOS API Function to do the job.

Run Time Statistics

Description

FreeRTOS can optionally collect information on the amount of processing time that has been used by each task. The `vTaskGetRunTimeStats()` API function can then be used to present this information in a tabular format, as shown on the right.

But first we should edit the implementation of the function ***uxTaskGetSystemState()*** Which this API depend on to collect statistics for the system, we edit it to collect the needed information from our EDF Ready List instead of the original Ready List in case of using the EDF scheduler as follows.

```

2720  #if (configUSE_EDF_SCHEDULER == 1)
2721
2722      uxTask += prvListTasksWithinSingleList( &(amp; pxTaskStatusArray[ uxTask ] ), &(xReadyTasksListEDF), eReady );
2723  #else
2724      do
2725      {
2726          uxQueue--;
2727          uxTask += prvListTasksWithinSingleList( &(amp; pxTaskStatusArray[ uxTask ] ), &(pxReadyTasksLists[ uxQueue ] ), eReady
2728      } while( uxQueue > ( UBaseType_t ) tskIDLE_PRIORITY ); /*lint !e961 MISRA exception as the casts are only redundant for
2729  #endif
2730

```

After then we configure our general purpose timer at suitable frequency 10 times the tick frequency and define other needed macros for this API to work.

Calling this API function and transmit the formatted ASCII information by UART and displaying them on the terminal we get the follows:

| | | |
|--------------------------------|-------|-----|
| I | | |
| DLE | 32624 | 35% |
| LOAD_2 | 11997 | 13% |
| UART | 115 | <1% |
| BUTTON_1_MONITOR | 29 | <1% |
| BUTTON_2_MONITOR | 28 | <1% |
| Transmitter | 26 | <1% |
| LOAD_1 | 46607 | 50% |
| This is 100 ms Periodic String | | |
| I | | |
| DLE | 34800 | 35% |
| LOAD_2 | 12758 | 13% |
| UART | 121 | <1% |
| BUTTON_1_MONITOR | 31 | <1% |
| BUTTON_2_MONITOR | 30 | <1% |
| Transmitter | 28 | <1% |
| LOAD_1 | 49652 | 50% |

IDLE task takes around 35% of the CPU time, the time that the CPU is free. That means the CPU load is around 65% from the six tasks.

Verifying our statistics with the Simso tool we get so close observations from ours.

| Observation Window: | | | |
|------------------------|------------|---------|--------------|
| from 0.00 to 200.00 ms | | | Configure... |
| | Total load | Payload | System load |
| CPU 1 | 0.6215 | 0.6215 | 0.0000 |
| Average | 0.6215 | 0.6215 | 0.0000 |

All these results prove a successful implementation of our EDF scheduler.

Thank you.